

# Turbo-charging the Sobel Operator

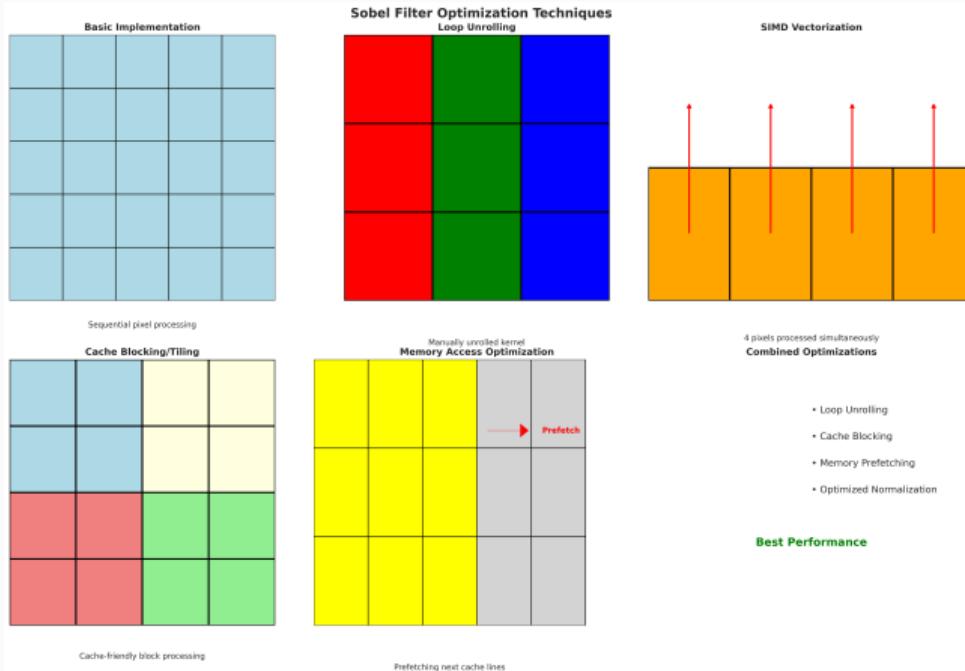
From NumPy to  $2200 \times$  with modern C++

---

Daniel Leon

June 6, 2025

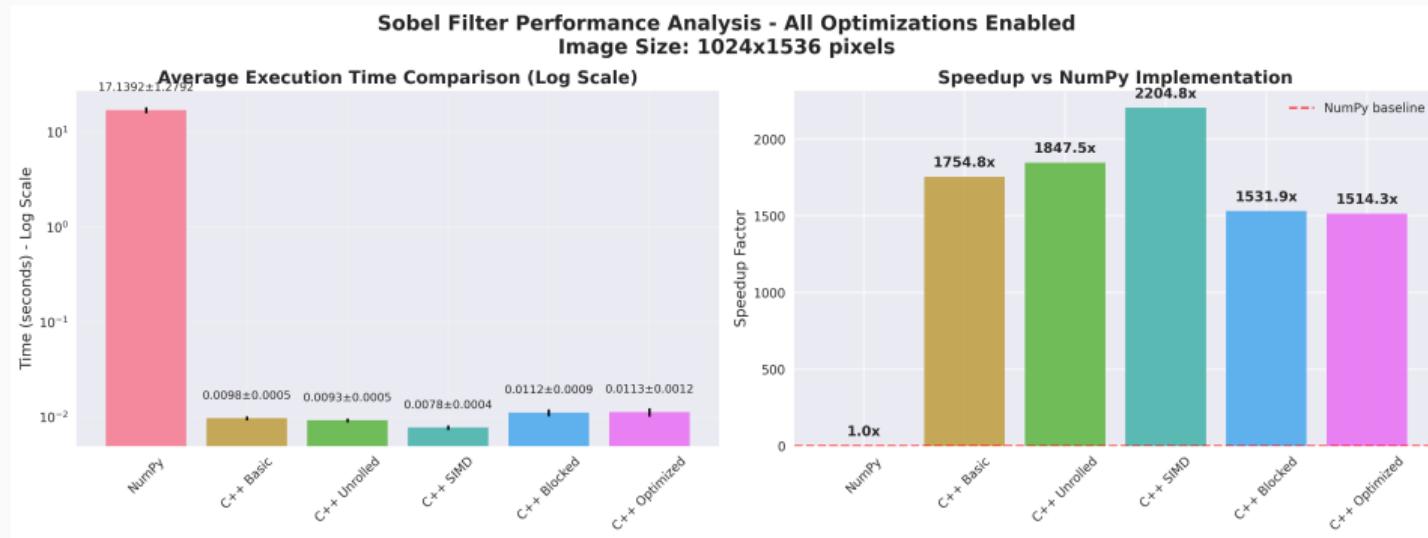
# The optimisation map



## Context

- Start: textbook kernel
- Add: loop unroll → SIMD → tiling
- 15-run mean on  $1024 \times 1536$  gray frame
- i9-10885H
- GCC 13
- pybind11

# Performance scoreboard



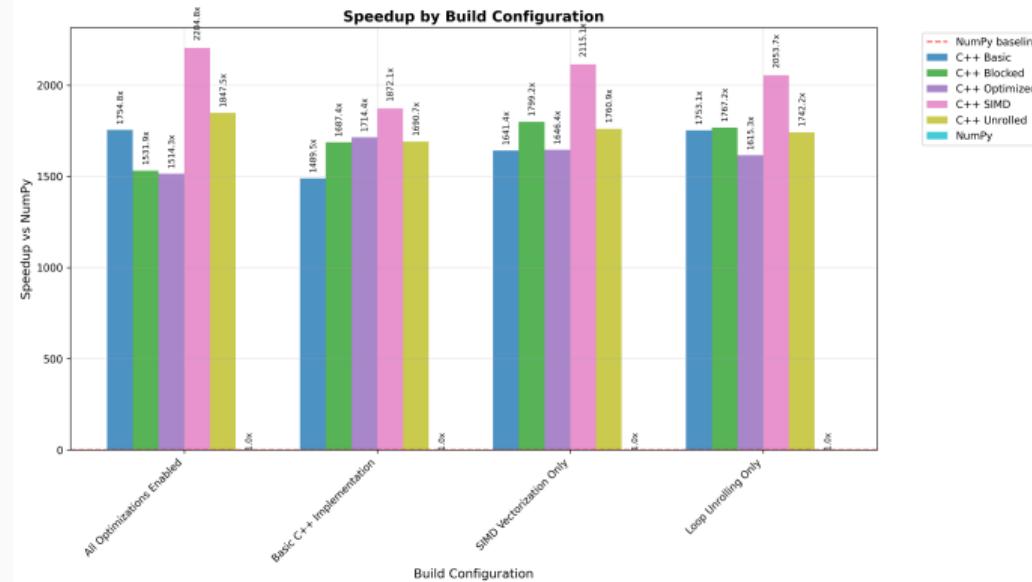
**2200× faster than NumPy**

# Do the edges still look right?



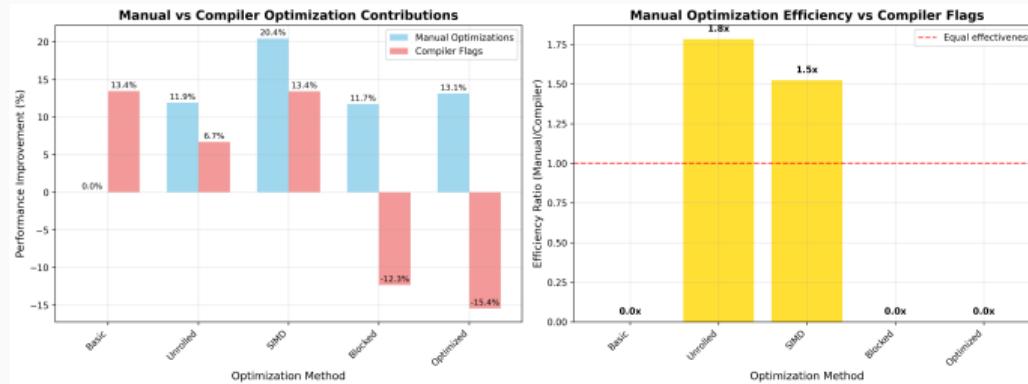
Bit-wise identical gradients — speed didn't cost accuracy.

# Which combo wins where?



SIMD + unroll dominates; blocking helps on big frames. Compiler flags alone leave **15 %** on the table.

# Manual vs compiler effort



- Manual tweaks still deliver **1.5–1.8x extra** over -O3 -ffast-math
- Best bang-for-buck: manual SIMD + -march=native (**+9 %**)
- Compiler flags plateau at 20 % — intrinsics still pay.

## System context scaling

### Test rig

- Intel i9-10885H @ 2.4 GHz
- 32 GB RAM, AVX2
- 8 cores / 16 threads

### Scaling hints

- Newer CPUs: +10–20 %
- Desktop clocks: +15–30 %
- AVX-512: up to 2× for SIMD path
- ARM M-series: NEON close to AVX2

Cache blocking gets sweeter with bigger L3 caches.

## Get the code

[github.com/dleon86/pybind11-vision-bench](https://github.com/dleon86/pybind11-vision-bench)

Runs on Linux macOS — PRs welcome for AVX-512 or GPU ports.